

Batch loading data and list columns

Daniel Anderson

Week 4

Agenda

- Discuss the midterm
 - Canvas quiz (10 points; please don't stress)
 - Take home (40 points)
- Review Lab 2
- `map_dfr` and batch-loading data

- Introduce list columns
- Contrast:
 - `group_by() %>% nest() %>% mutate() %>% map()` with
 - `nest_by() %>% summarize()`
- In-class midterm (last 25 minutes)

Review Lab 2

Learning objectives

- Understand when `map_dfr` can and should be applied
- Better understand file paths, and how `{fs}` can help
- Be able to batch load data of a specific type within a mixed-type directory
- Use filenames to pull data

Learning objectives (cont.)

- Understand list columns and how they relate to `base::split`
- Fluently nest/unnest data frames
- Understand why `tidyr::nest` can be a powerful framework (data frames) and when `tidyr::unnest` can/should be used to move out of nested data frames and into a standard data frame.

Midterm

Questions?

Let's look at the take-home portion

map_dfr

- If each iteration returns a data frame, you can use `map_dfr` to automatically bind all the data frames together.

Example

- Create a function that simulates data (please copy the code and follow along)

```
set.seed(8675309)
simulate <- function(n, mean = 0, sd = 1) {
  tibble(sample_id = seq_len(n),
          sample = rnorm(n, mean, sd))
}
simulate(10)
```

```
## # A tibble: 10 × 2
##   sample_id      sample
##   <int>         <dbl>
## 1         1 -0.9965824
## 2         2  0.7218241
## 3         3 -0.6172088
## 4         4  2.029392
## 5         5  1.065416
## 6         6  0.9872197
## 7         7  0.02745393
## 8         8  0.6728723
## 9         9  0.5720665
## 10        10  0.9036777
```

Two more quick examples

```
simulate(3, 100, 10)
```

```
## # A tibble: 3 × 2
##   sample_id    sample
##   <int>      <dbl>
## 1         1  84.50448
## 2         2 110.2264
## 3         3 101.5008
```

```
simulate(5, -10, 1.5)
```

```
## # A tibble: 5 × 2
##   sample_id    sample
##   <int>      <dbl>
## 1         1 -10.98995
## 2         2 -11.49188
## 3         3  -7.041312
## 4         4 -10.66270
## 5         5 -11.35096
```

Simulation

- Assume we want to vary the sample size from 10 to 150 by increments of 5
- `mean` stays constant at 100, `sd` is constant at 10

Try with `purrr::map`

02:00

```
library(tidyverse)
sims <- map(seq(10, 150, 5), simulate, mean = 100, sd = 10)
```

```
sims[1]
```

```
## [[1]]
## # A tibble: 10 × 2
##   sample_id sample
##   <int>     <dbl>
## 1         1 103.7618
## 2         2 111.5353
## 3         3 115.7490
## 4         4 105.8853
## 5         5  93.84955
## 6         6  97.71089
## 7         7 100.6392
## 8         8  96.86526
## 9         9  97.51501
## 10        10  98.46205
```

```
sims[2]
```

```
## [[1]]
## # A tibble: 15 × 2
##   sample_id sample
##   <int>     <dbl>
## 1         1  93.64743
## 2         2  99.96206
## 3         3 100.4562
## 4         4 106.8407
## 5         5  97.47957
## 6         6  98.48961
## 7         7  91.25069
## 8         8  80.23099
## 9         9 102.3766
## 10        10 100.3609
## 11        11 101.3490
## 12        12 101.1758
## 13        13  91.74411
## 14        14  78.64764
## 15        15 102.1421
```

Swap for `map_dfr`

Try it – what happens?

```
sims_df <- map_dfr(seq(10, 150, 5), simulate, 100, 10)
sims_df
```

```
## # A tibble: 2,320 × 2
##   sample_id    sample
##   <int>      <dbl>
## 1         1  85.64361
## 2         2 103.6789
## 3         3  94.71782
## 4         4 103.1350
## 5         5  99.78701
## 6         6 105.3462
## 7         7 100.0653
## 8         8  94.28314
## 9         9 108.8872
## 10        10 106.0850
## # ... with 2,310 more rows
```

01:00

Notice a problem here

```
sims_df[1:15, ]
```

```
## # A tibble: 15 × 2
##   sample_id sample
##   <int>     <dbl>
## 1         1  85.64361
## 2         2 103.6789
## 3         3  94.71782
## 4         4 103.1350
## 5         5  99.78701
## 6         6 105.3462
## 7         7 100.0653
## 8         8  94.28314
## 9         9 108.8872
## 10        10 106.0850
## 11         1  89.49968
## 12         2  86.99898
## 13         3  85.38054
## 14         4  99.10690
## 15         5 105.0088
```

.id argument

```
sims_df2 <- map_dfr(seq(10, 150, 5), simulate, 100, 10,  
                   .id = "iteration")  
sims_df2[1:15, ]
```

```
## # A tibble: 15 × 3  
##   iteration sample_id    sample  
##   <chr>      <int>      <dbl>  
## 1 1          1 112.1250  
## 2 1          2  88.07056  
## 3 1          3 108.3908  
## 4 1          4 100.8193  
## 5 1          5 102.1545  
## 6 1          6 113.5398  
## 7 1          7 101.4171  
## 8 1          8  99.33668  
## 9 1          9 100.2855  
## 10 1         10  90.22043  
## 11 2          1  91.08882  
## 12 2          2 107.3664  
## 13 2          3 101.1745  
## 14 2          4  96.82053  
## 15 2          5 105.5844
```

.id: Either a string or NULL. If a string, the output will contain a variable with that name, storing either the name (if `.x` is named) or the index (if `.x` is unnamed) of the input. If NULL, the default, no variable will be created.

– `{purrr}` documentation

setNames

```
sample_size <- seq(10, 150, 5)
sample_size
```

```
## [1] 10 15 20 25 30 35 40 45 50 55 60 65 70 75 80 85 90
## [22] 115 120 125 130 135 140 145 150
```

```
sample_size <- setNames(sample_size,
                        english::english(seq(10, 150, 5)))
sample_size[1:15]
```

```
##          ten          fifteen          twenty  twenty-five          thirty  thirty
##          10           15           20           25           30           30
##       forty  forty-five          fifty  fifty-five          sixty  sixty
##          40           45           50           55           60           60
##       seventy seventy-five          eighty
##          70           75           80
```

Try again

```
sims_df3 <- map_dfr(sample_size, simulate, 100, 10,  
                    .id = "n")  
sims_df3[1:15, ]
```

```
## # A tibble: 15 × 3  
##       n      sample_id      sample  
##   <chr>      <int>      <dbl>  
## 1 ten           1  98.94914  
## 2 ten           2 101.6824  
## 3 ten           3  88.16447  
## 4 ten           4  90.13604  
## 5 ten           5  85.53591  
## 6 ten           6  90.69977  
## 7 ten           7 105.8858  
## 8 ten           8  89.12978  
## 9 ten           9 114.4982  
## 10 ten          10 111.6440  
## 11 fifteen      1 103.2732  
## 12 fifteen      2 106.8949  
## 13 fifteen      3  88.83591  
## 14 fifteen      4 105.5402  
## 15 fifteen      5 112.6581
```

Another quick example

broom::tidy

- The {broom} package helps us extract model output in a tidy format

```
lm(tvhours ~ age, gss_cat) %>%  
  broom::tidy()
```

```
## # A tibble: 2 × 5  
##   term          estimate  std.error statistic    p.value  
##   <chr>          <dbl>      <dbl>    <dbl>    <dbl>  
## 1 (Intercept)  1.992391    0.06980966  28.54033 4.672161e-173  
## 2 age          0.02095679  0.001387361  15.10551 4.689310e- 51
```

Fit separate models by year

Again – probs not best statistically

```
split(gss_cat, gss_cat$year) %>%  
  map_dfr(  
    ~lm(tvhours ~ age, .x) %>%  
      broom::tidy()  
  )
```

```
## # A tibble: 16 × 5  
##   term          estimate std.error statistic    p.value  
##   <chr>          <dbl>      <dbl>    <dbl>    <dbl>  
## 1 (Intercept)  2.080163  0.1709061  12.17138  7.995632e-33  
## 2 age          0.01948584 0.003485199  5.591027  2.599011e- 8  
## 3 (Intercept)  2.078999  0.2176829   9.550583  1.191266e-20  
## 4 age          0.01963575 0.004400292  4.462375  9.137366e- 6  
## 5 (Intercept)  1.767990  0.2464509   7.173804  1.531756e-12  
## 6 age          0.02386070 0.005031548  4.742218  2.459650e- 6  
## 7 (Intercept)  2.096054  0.1496431  14.00702  1.419772e-42  
## 8 age          0.01781388 0.002977289  5.983256  2.589482e- 9  
## 9 (Intercept)  1.855278  0.2156381   8.603668  2.167351e-17  
## 10 age         0.02390720 0.004314567  5.541043  3.628675e- 8  
## 11 (Intercept)  2.068914  0.2096397   9.868903  2.896085e-22  
## 12 age         0.01989505 0.004086638  4.868317  1.251234e- 6  
## 13 (Intercept)  1.878070  0.2258400   8.315932  2.280108e-16
```

.id

In cases like the preceding, `.id` becomes invaluable

```
split(gss_cat, gss_cat$year) %>%  
  map_dfr(  
    ~lm(tvhours ~ age, .x) %>%  
      broom::tidy(),  
    .id = "year"  
  )
```

```
## # A tibble: 16 × 6  
##   year term      estimate std.error statistic    p.value  
##   <chr> <chr>      <dbl>      <dbl>      <dbl>    <dbl>  
## 1 2000 (Intercept) 2.080163 0.1709061 12.17138 7.995632e-33  
## 2 2000 age      0.01948584 0.003485199 5.591027 2.599011e- 8  
## 3 2002 (Intercept) 2.078999 0.2176829 9.550583 1.191266e-20  
## 4 2002 age      0.01963575 0.004400292 4.462375 9.137366e- 6  
## 5 2004 (Intercept) 1.767990 0.2464509 7.173804 1.531756e-12  
## 6 2004 age      0.02386070 0.005031548 4.742218 2.459650e- 6  
## 7 2006 (Intercept) 2.096054 0.1496431 14.00702 1.419772e-42  
## 8 2006 age      0.01781388 0.002977289 5.983256 2.589482e- 9  
## 9 2008 (Intercept) 1.855278 0.2156381 8.603668 2.167351e-17  
## 10 2008 age      0.02390720 0.004314567 5.541043 3.628675e- 8  
## 11 2010 (Intercept) 2.068914 0.2096397 9.868903 2.896085e-22  
## 12 2010 age      0.01989505 0.004086638 4.868317 1.251234e-26 / 107
```

Break

05:00

Batch— loading data

Please follow along

{fs}

Could we apply `map_dfr` here?

```
# install.packages("fs")  
library(fs)  
dir_ls(here::here("data"))
```

```
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p  
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p
```

```
dir_ls(here::here("data", "pfiles_sim"))
```

```
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p  
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p  
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p  
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p  
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p  
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p  
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p  
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p  
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p  
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p  
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p
```


Limit files

- We really only want the `.CSV`
 - That happens to be the only thing that's in there but that's regularly not the case

```
dir_ls(here::here("data", "pfiles_sim"), glob = "*.csv")
```

[illegible]

Batch load

- Loop through the directories and `import` or `read_csv`

```
files <- dir_ls(  
  here::here("data", "pfiles_sim"),  
  glob = "*.csv"  
)  
batch <- map_dfr(files, read_csv)  
batch
```

```
## # A tibble: 15,945 × 22  
##   Entry   Theta Status Count RawScore      SE Infit Infit_Z Outfit Outfit_Z  
##   <dbl>   <dbl>   <dbl> <dbl>   <dbl>   <dbl> <dbl>   <dbl>   <dbl>   <dbl>  
## 1    123  1.2687     1    36      23 0.3713  0.93   -0.34    0.82   -0.34  
## 2     88  1.5541     1    36      25 0.3852  0.95   -0.37    0.81   -0.37  
## 3    105  3.2773     1    36      33 0.6187  0.9   -0.04    1.63    1.63  
## 4    153  4.4752     1    36      35 1.0234  0.93    0.23    0.35   -0.35  
## 5    437  2.6655     1    36      31 0.5008  0.92   -0.18    0.88   -0.18  
## 6    307  5.7137     0    36      36 1.8371  1      0      1      0  
## 7    305  3.7326     1    36      34 0.7408  1.06    0.31    0.86    0.86  
## 8     42  0.609      1    36      18 0.36    1.55    2.56    1.74    3.14  
## 9     59 -2.623      1    36       3 1.0344  0.85    0.06    0.17   -0.17  
## 10   304  5.7137     0    36      36 1.8371  1      0      1      0  
## # ... with 15,935 more rows, and 11 more variables: PointMeasureCorr <dbl>,  
## #   ObservMatch <dbl>, ExpectMatch <dbl>, PointMeasureExpected <dbl>, RMSE <dbl>
```

Problem

- We've lost a lot of info – no way to identify which file is which

Try to fix it!

02:00

Add id

```
batch2 <- map_dfr(files, read_csv, .id = "file")  
batch2
```

```
## # A tibble: 15,945 × 23  
## # ... with 15,935 more rows, and 23 more variables: file <chr>, Entry <dbl>  
## #   Status <dbl>, Count <dbl>, RawScore <dbl>, SE <dbl>, Infit <dbl>, In  
## #   Outfit <dbl>, Outfit_Z <dbl>, Displacement <dbl>, PointMeasureCorr <  
## #   Weight <dbl>, ObservMatch <dbl>, ExpectMatch <dbl>, PointMeasureExpe  
## #   RMSR <dbl>, WMLE <dbl>, testeventid <dbl>, ssid <dbl>, asmtprmrydsbl  
## #   asmtscndrydsbltycd <dbl>
```

Note – the `file` column contains the full path, which is so long it makes no rows print

```
batch2 %>%  
  count(file)
```

```
## # A tibble: 31 × 2  
## # ... with 21 more rows, and 2 more variables: file <chr>, n <int>
```

- Still not terrifically useful. What can we do?

Step 1

- Remove the `here::here` path from string

```
batch2 <- batch2 %>%  
  mutate(  
    file = str_replace_all(  
      file,  
      here::here("data", "pfiles_sim"),  
      ""  
    )  
  )  
  
batch2 %>%  
  count(file)
```

```
## # A tibble: 31 × 2  
##   file                                n  
##   <chr>                             <int>  
## 1 /g11ELApfiles18_sim.csv           453  
## 2 /g11Mathpfiles18_sim.csv          460  
## 3 /g11Rdgpfiles18_sim.csv           453  
## 4 /g11Sciencepfiles18_sim.csv       438  
## 5 /g11Wripfiles18_sim.csv           453  
## 6 /g3ELApfiles18_sim.csv           540
```

Pull out pieces you need

- Regular expressions are most powerful here
 - We haven't talked about them much
- Try RegExplain

Pull grade

- Note – I'm not expecting you to just suddenly be able to do this. This is more for illustration. There's also other ways you could extract the same info

```
batch2 %>%  
  mutate(  
    grade = str_replace_all(  
      file,  
      "/g(\\d?\\d).+",  
      "\\1"  
    )  
  ) %>%  
  select(file, grade)
```

```
## # A tibble: 15,945 × 2  
##   file                                grade  
##   <chr>                             <chr>  
## 1 /g11ELApfiles18_sim.csv 11  
## 2 /g11ELApfiles18_sim.csv 11  
## 3 /g11ELApfiles18_sim.csv 11  
## 4 /g11ELApfiles18_sim.csv 11  
## 5 /g11ELApfiles18_sim.csv 11
```


parse_number

- In this case `parse_number` also works – but note that it would not work to extract the year

```
batch2 %>%  
  mutate(grade = parse_number(file)) %>%  
  select(file, grade)
```

```
## # A tibble: 15,945 × 2  
##   file                                grade  
##   <chr>                             <dbl>  
## 1 /g11ELApfiles18_sim.csv           11  
## 2 /g11ELApfiles18_sim.csv           11  
## 3 /g11ELApfiles18_sim.csv           11  
## 4 /g11ELApfiles18_sim.csv           11  
## 5 /g11ELApfiles18_sim.csv           11  
## 6 /g11ELApfiles18_sim.csv           11  
## 7 /g11ELApfiles18_sim.csv           11  
## 8 /g11ELApfiles18_sim.csv           11  
## 9 /g11ELApfiles18_sim.csv           11  
## 10 /g11ELApfiles18_sim.csv          11  
## # ... with 15,935 more rows
```

Extract year

```
batch2 %>%  
  mutate(  
    grade = str_replace_all(  
      file, "/g(\\d?\\d).+", "\\1"  
    ),  
    year = str_replace_all(  
      file, ".+files(\\d\\d)_sim.", "\\1"  
    )  
  ) %>%  
  select(file, grade, year)
```

```
## # A tibble: 15,945 × 3  
##   file                                grade year  
##   <chr>                             <chr> <chr>  
## 1 /g11ELApfiles18_sim.csv 11      18  
## 2 /g11ELApfiles18_sim.csv 11      18  
## 3 /g11ELApfiles18_sim.csv 11      18  
## 4 /g11ELApfiles18_sim.csv 11      18  
## 5 /g11ELApfiles18_sim.csv 11      18  
## 6 /g11ELApfiles18_sim.csv 11      18  
## 7 /g11ELApfiles18_sim.csv 11      18  
## 8 /g11ELApfiles18_sim.csv 11      18  
## 9 /g11ELApfiles18_sim.csv 11      18  
## 10 /g11ELApfiles18_sim.csv 11      18
```

Extract Content Area

```
batch2 %>%  
  mutate(grade = str_replace_all(file,  
                                   "/g(\\d?\\d).+",  
                                   "\\1"),  
         year = str_replace_all(file,  
                                 "+files(\\d\\d)_sim.",  
                                 "\\1"),  
         content = str_replace_all(file,  
                                    "/g\\d?\\d(\\d+)pfiles.",  
                                    "\\1")) %>%  
  select(file, grade, year, content)
```

```
## # A tibble: 15,945 × 4  
##   file                                grade year content  
##   <chr>                             <chr> <chr> <chr>  
## 1 /g11ELApfiles18_sim.csv          11    18   ELA  
## 2 /g11ELApfiles18_sim.csv          11    18   ELA  
## 3 /g11ELApfiles18_sim.csv          11    18   ELA  
## 4 /g11ELApfiles18_sim.csv          11    18   ELA  
## 5 /g11ELApfiles18_sim.csv          11    18   ELA  
## 6 /g11ELApfiles18_sim.csv          11    18   ELA  
## 7 /g11ELApfiles18_sim.csv          11    18   ELA  
## 8 /g11ELApfiles18_sim.csv          11    18   ELA  
## 9 /g11ELApfiles18_sim.csv          11    18   ELA
```

Double checks: grade

```
batch2 %>%
  mutate(grade = str_replace_all(file,
                                   "/g(\\d?\\d).+",
                                   "\\1"),
         year = str_replace_all(file,
                                   ".+files(\\d\\d)_sim.+",
                                   "\\1"),
         content = str_replace_all(file,
                                   "/g\\d?\\d(\\d.+?)pfiles.+",
                                   "\\1")) %>%
  select(file, grade, year, content) %>%
  count(grade)
```

```
## # A tibble: 7 × 2
##   grade      n
##   <chr> <int>
## 1 11      2257
## 2 3       2156
## 3 4       2341
## 4 5       2632
## 5 6       2216
## 6 7       1962
## 7 8       2381
```

Double checks: year

```
batch2 %>%
  mutate(grade = str_replace_all(file,
                                   "/g(\\d?\\d).+",
                                   "\\1"),
         year = str_replace_all(file,
                                   ".+files(\\d\\d)_sim.",
                                   "\\1"),
         content = str_replace_all(file,
                                   "/g\\d?\\d(\\d.+?)pfiles.",
                                   "\\1")) %>%
  select(file, grade, year, content) %>%
  count(year)
```

```
## # A tibble: 1 × 2
##   year      n
##   <chr> <int>
## 1 18    15945
```

Double checks: content

```
batch2 %>%
  mutate(grade = str_replace_all(file,
                                   "/g(\\d?\\d).+",
                                   "\\1"),
         year = str_replace_all(file,
                                   ".+files(\\d\\d)_sim.+",
                                   "\\1"),
         content = str_replace_all(file,
                                   "/g\\d?\\d(\\d.+?)pfiles.+",
                                   "\\1")) %>%
  select(file, grade, year, content) %>%
  count(content)
```

```
## # A tibble: 5 × 2
##   content      n
##   <chr>    <int>
## 1 ELA      3627
## 2 Math     3629
## 3 Rdg      3627
## 4 Science  1435
## 5 Wri      3627
```

Finalize

```
d <- batch2 %>%
  mutate(
    grade = str_replace_all(
      file, "/g(\\d?\\d).+", "\\1"
    ),
    grade = as.integer(grade),
    year = str_replace_all(
      file, ".+files(\\d\\d)_sim.+", "\\1"
    ),
    year = as.integer(year),
    content = str_replace_all(
      file, "/g\\d?\\d(.+)pfiles.+", "\\1"
    )
  ) %>%
  select(-file) %>%
  select(
    ssid, grade, year, content, testeventid,
    asmtprmrydsbltycd, asmtscndrydsbltycd, Entry:WMLE
  )
```


Final product

- In this case, we basically have a tidy data frame already!
- We've reduced our problem from 31 files to a single file

d

```
## # A tibble: 15,945 × 25
##       ssid grade  year content testeventid asmtprmrtdsbltycd asmtscndry
##       <dbl> <int> <int> <chr>          <dbl>          <dbl>
##  1  9466908    11    11 ELA             148933            0
##  2  7683685    11    11 ELA             147875           10
##  3  9025693    11    11 ELA             143699           40
##  4 10099824    11    11 ELA             143962           82
##  5 18886078    11    11 ELA             150680           10
##  6 10606750    11    11 ELA             144583           80
##  7 10541306    11    11 ELA             145204           50
##  8  7632967    11    11 ELA             148926           10
##  9  7661118    11    11 ELA             148893           50
## 10 10547177    11    11 ELA             144583           82
## # ... with 15,935 more rows, and 17 more variables: Theta <dbl>, Status <dbl>,
## #   Count <dbl>, RawScore <dbl>, SE <dbl>, Infit <dbl>, Infit_Z <dbl>,
## #   Outfit_Z <dbl>, Displacement <dbl>, PointMeasureCorr <dbl>, Weight <dbl>,
## #   ObservMatch <dbl>, ExpectMatch <dbl>, PointMeasureExpected <dbl>, RMSE <dbl>
```


Quick look at distributions



Summary stats

```
d %>%
  group_by(grade, content, asmtprmrydsbltycd) %>%
  summarize(mean = mean(Theta)) %>%
  pivot_wider(names_from = content,
              values_from = mean)
```

```
## # A tibble: 77 × 7
## # Groups:   grade [7]
##   grade asmtprmrydsbltycd      ELA      Math      Rdg      Wri
##   <int>          <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1      3              0 -0.07361 -1.21055  1.010455  1.612308
## 2      3             10  0.3700416 -0.8182091  0.5184354  0.3206475
## 3      3             20 -0.06335  -1.2514   1.52      -0.5775
## 4      3             40 -1.877683  -3.56365  -1.761667 -0.7514286
## 5      3             50  0.9462857 -0.09186957  0.9791176  1.191481
## 6      3             60  0.840775   1.040375   2.181111  1.067
## 7      3             70 -1.104049  -1.517955  -0.8454839 -1.005625
## 8      3             74  0.996      0.0208375  0.6       1.2925
## 9      3             80 -0.144304  -0.5325596  0.6791667  0.2686301
## 10     3             82  0.3708244 -1.080988  0.5676650  0.3440741
## # ... with 67 more rows
```

Backing up a bit

- What if we wanted only math files?

```
dir_ls(here::here("data", "pfiles_sim"), regexp = "Math")
```

```
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p
## /Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/data/p
```

Only Grade 5

You try

```
g5_paths <- dir_ls(  
  here::here("data", "pfiles_sim"),  
  regexp = "g5"  
)
```

03:00

The rest is the same

```
g5 <- map_dfr(g5_paths, read_csv, .id = "file") %>%  
  mutate(  
    file = str_replace_all(  
      file,  
      here::here("data", "pfiles_sim"),  
      ""  
    )  
  )  
g5
```

```
## # A tibble: 2,632 × 23
```

| ## | file | Entry | Theta | Status | Count | RawScore | SE | Inf |
|----|---------------------------|-------|---------|--------|-------|----------|--------|-------|
| ## | <chr> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> |
| ## | 1 /g5ELApfiles18_sim.csv | 375 | 3.154 | 1 | 36 | 32 | 0.551 | 1. |
| ## | 2 /g5ELApfiles18_sim.csv | 305 | 0.3662 | 1 | 36 | 16 | 0.3894 | 0. |
| ## | 3 /g5ELApfiles18_sim.csv | 163 | -4.9547 | -1 | 36 | 0 | 1.8495 | 1 |
| ## | 4 /g5ELApfiles18_sim.csv | 524 | -4.9547 | -1 | 36 | 0 | 1.8495 | 1 |
| ## | 5 /g5ELApfiles18_sim.csv | 81 | 3.154 | 1 | 36 | 32 | 0.551 | 0. |
| ## | 6 /g5ELApfiles18_sim.csv | 325 | 1.7156 | 1 | 36 | 25 | 0.3997 | 1. |
| ## | 7 /g5ELApfiles18_sim.csv | 163 | 1.8786 | 1 | 36 | 26 | 0.4078 | 0. |
| ## | 8 /g5ELApfiles18_sim.csv | 116 | 5.9323 | 0 | 36 | 36 | 1.8373 | 1 |
| ## | 9 /g5ELApfiles18_sim.csv | 273 | 1.4052 | 1 | 36 | 23 | 0.3891 | 1. |
| ## | 10 /g5ELApfiles18_sim.csv | 202 | 1.8786 | 1 | 36 | 26 | 0.4078 | 0. |

```
## # ... with 2,622 more rows, and 14 more variables: Outfit <dbl>, Outfit_Z
```

Base equivalents

```
list.files(here::here("data", "pfiles_sim"))
```

```
## [1] "g11ELApfiles18_sim.csv"      "g11Mathpfiles18_sim.csv"
## [3] "g11Rdgpfiles18_sim.csv"      "g11Sciencepfiles18_sim.csv"
## [5] "g11Wripfiles18_sim.csv"      "g3ELApfiles18_sim.csv"
## [7] "g3Mathpfiles18_sim.csv"      "g3Rdgpfiles18_sim.csv"
## [9] "g3Wripfiles18_sim.csv"      "g4ELApfiles18_sim.csv"
## [11] "g4Mathpfiles18_sim.csv"      "g4Rdgpfiles18_sim.csv"
## [13] "g4Wripfiles18_sim.csv"      "g5ELApfiles18_sim.csv"
## [15] "g5Mathpfiles18_sim.csv"      "g5Rdgpfiles18_sim.csv"
## [17] "g5Sciencepfiles18_sim.csv"   "g5Wripfiles18_sim.csv"
## [19] "g6ELApfiles18_sim.csv"      "g6Mathpfiles18_sim.csv"
## [21] "g6Rdgpfiles18_sim.csv"      "g6Wripfiles18_sim.csv"
## [23] "g7ELApfiles18_sim.csv"      "g7Mathpfiles18_sim.csv"
## [25] "g7Rdgpfiles18_sim.csv"      "g7Wripfiles18_sim.csv"
## [27] "g8ELApfiles18_sim.csv"      "g8Mathpfiles18_sim.csv"
## [29] "g8Rdgpfiles18_sim.csv"      "g8Sciencepfiles18_sim.csv"
## [31] "g8Wripfiles18_sim.csv"
```

Full path

```
list.files(here::here("data", "pfiles_sim"), full.names = TRUE)
```

```
## [1] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [2] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [3] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [4] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [5] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [6] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [7] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [8] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [9] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [10] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [11] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [12] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [13] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [14] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [15] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [16] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [17] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [18] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [19] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [20] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [21] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [22] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [23] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
## [24] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/"
```

Only csvs

```
list.files(here::here("data", "pfiles_sim"),  
          full.names = TRUE,  
          pattern = "*.csv")
```

```
## [1] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [2] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [3] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [4] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [5] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [6] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [7] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [8] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [9] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [10] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [11] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [12] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [13] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [14] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [15] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [16] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [17] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [18] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [19] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [20] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/  
## [21] "/Users/daniel/Teaching/data_sci_specialization/2021-22/c3-fp-2022/
```


Why not use base?

- We could, but `{fs}` plays a little nicer with `{purrr}`

```
files <- list.files(  
  here::here("data", "pfiles_sim"),  
  pattern = "*.csv"  
)  
batch3 <- map_dfr(files, read_csv, .id = "file")
```

```
## Error: 'g11ELApfiles18_sim.csv' does not exist in current working direct
```

- Need to return full names

```
files
```

```
## [1] "g11ELApfiles18_sim.csv"      "g11Mathpfiles18_sim.csv"  
## [3] "g11Rdgpfiles18_sim.csv"     "g11Sciencepfiles18_sim.csv"  
## [5] "g11Wripfiles18_sim.csv"     "g3ELApfiles18_sim.csv"  
## [7] "g3Mathpfiles18_sim.csv"     "g3Rdgpfiles18_sim.csv"  
## [9] "g3Wripfiles18_sim.csv"     "g4ELApfiles18_sim.csv"  
## [11] "g4Mathpfiles18_sim.csv"     "g4Rdgpfiles18_sim.csv"  
## [13] "g4Wripfiles18_sim.csv"     "g5ELApfiles18_sim.csv"
```

Try again

```
files <- list.files(here::here("data", "pfiles_sim"),
                    pattern = "*.csv",
                    full.names = TRUE)
batch3 <- map_dfr(files, read_csv, .id = "file")
batch3
```

```
## # A tibble: 15,945 × 23
##   file Entry  Theta Status Count RawScore      SE Infit Infit_Z Outfit
##   <chr> <dbl>   <dbl>   <dbl> <dbl>   <dbl>   <dbl> <dbl>   <dbl>   <dbl>
## 1 1      123  1.2687     1     36     23 0.3713  0.93   -0.34  0.82
## 2 1      88  1.5541     1     36     25 0.3852  0.95   -0.37  0.81
## 3 1     105  3.2773     1     36     33 0.6187  0.9    -0.04  1.63
## 4 1     153  4.4752     1     36     35 1.0234  0.93    0.23  0.35
## 5 1     437  2.6655     1     36     31 0.5008  0.92   -0.18  0.88
## 6 1     307  5.7137     0     36     36 1.8371  1       0     1
## 7 1     305  3.7326     1     36     34 0.7408  1.06    0.31  0.86
## 8 1      42  0.609      1     36     18 0.36    1.55    2.56  1.74
## 9 1      59 -2.623      1     36      3 1.0344  0.85    0.06  0.17
## 10 1     304  5.7137     0     36     36 1.8371  1       0     1
## # ... with 15,935 more rows, and 12 more variables: Displacement <dbl>,
## #   PointMeasureCorr <dbl>, Weight <dbl>, ObservMatch <dbl>, ExpectMatch
## #   PointMeasureExpected <dbl>, RMSR <dbl>, WMLE <dbl>, testeventid <dbl>,
## #   asmtprmrydsbltycd <dbl>, asmtscndrydsbltycd <dbl>
```

indexes

- The prior example gave us indexes, rather than the file path. Why?

No names

```
names(files)
```

```
## NULL
```

- We **need** the file path! An index isn't nearly as useful.

Base method that works

```
files <- list.files(here::here("data", "pfiles_sim"),
                   pattern = "*.csv",
                   full.names = TRUE)
files <- setNames(files, files)

batch4 <- map_dfr(files, read_csv, .id = "file")
batch4
```

```
## # A tibble: 15,945 × 23
## # ... with 15,935 more rows, and 23 more variables: file <chr>, Entry <dbl>,
## #   Status <dbl>, Count <dbl>, RawScore <dbl>, SE <dbl>, Infit <dbl>, In
## #   Outfit <dbl>, Outfit_Z <dbl>, Displacement <dbl>, PointMeasureCorr <
## #   Weight <dbl>, ObservMatch <dbl>, ExpectMatch <dbl>, PointMeasureExpe
## #   RMSR <dbl>, WMLE <dbl>, testeventid <dbl>, ssid <dbl>, asmtprmrydsbl
## #   asmtscndrydsbltycd <dbl>
```

My recommendation

- If you're working interactively, no reason not to use `{fs}`
- If you are building **functions** that take paths, might be worth considering skipping the dependency

Note

I am **not** saying skip it, but rather that you should **consider** whether it is really needed or not.

List columns

Comparing models

Let's say we wanted to fit/compare a set of models for each content area

1. `lm(Theta ~ asmtprmrydsbltycd)`
2. `lm(Theta ~ asmtprmrydsbltycd +
asmtscndrydsbltycd)`
3. `lm(Theta ~ asmtprmrydsbltycd *
asmtscndrydsbltycd)`

Data pre-processing

- The disability variables are stored as numbers, we need them as factors
- We'll make the names easier in the process

```
d <- d %>%  
  mutate(primary = as.factor(asmtprmrydsbltycd),  
         secondary = as.factor(asmtscndrydsbltycd))
```

If you're interested in what the specific codes refer to, see [here](#).

Split the data

The base method we've been using...

```
splt_content <- split(d, d$content)
str(splt_content)
```

```
## List of 5
## $ ELA      : tibble [3,627 × 27] (S3: tbl_df/tbl/data.frame)
## ..$ ssid      : num [1:3627] 9466908 7683685 9025693 1009982
## ..$ grade     : int [1:3627] 11 11 11 11 11 11 11 11 11 11
## ..$ year      : int [1:3627] 11 11 11 11 11 11 11 11 11 11
## ..$ content    : chr [1:3627] "ELA" "ELA" "ELA" "ELA" ...
## ..$ testeventid : num [1:3627] 148933 147875 143699 143962 150
## ..$ asmtprmrydsblycd : num [1:3627] 0 10 40 82 10 80 50 10 50 82
## ..$ asmtscndrydsblycd : num [1:3627] 0 0 20 0 0 80 0 0 0 0
## ..$ Entry      : num [1:3627] 123 88 105 153 437 307 305 42 5
## ..$ Theta      : num [1:3627] 1.27 1.55 3.28 4.48 2.67 ...
## ..$ Status     : num [1:3627] 1 1 1 1 1 0 1 1 1 0
## ..$ Count      : num [1:3627] 36 36 36 36 36 36 36 36 36 36
## ..$ RawScore   : num [1:3627] 23 25 33 35 31 36 34 18 3 36
## ..$ SE         : num [1:3627] 0.371 0.385 0.619 1.023 0.501
## ..$ Infit      : num [1:3627] 0.93 0.95 0.9 0.93 0.92 1 1.06
## ..$ Infit_Z    : num [1:3627] -0.34 -0.37 -0.04 0.23 -0.18 0
## ..$ Outfit     : num [1:3627] 0.82 0.81 1.63 0.35 0.88 1 0.86
## ..$ Outfit_Z   : num [1:3627] -0.62 -0.56 1.03 -0.16 -0.12 0
```

We could use this method

```
m1 <- map(
  splt_content,
  ~lm(Theta ~ asmtprmrydsbltycd, data = .x)
)

m2 <- map(
  splt_content,
  ~lm(Theta ~ asmtprmrydsbltycd + asmtscndrydsbltycd,
      data = .x)
)

m3 <- map(
  splt_content,
  ~lm(Theta ~ asmtprmrydsbltycd * asmtscndrydsbltycd,
      data = .x)
)
```

- Then conduct tests to see which model fit better, etc.

Alternative

- Create a data frame with a list column

```
by_content <- d %>%  
  group_by(content) %>%  
  nest()  
by_content
```

```
## # A tibble: 5 × 2  
## # Groups:   content [5]  
##   content data  
##   <chr>    <list>  
## 1 ELA      <tibble [3,627 × 26]>  
## 2 Math     <tibble [3,629 × 26]>  
## 3 Rdg      <tibble [3,627 × 26]>  
## 4 Science  <tibble [1,435 × 26]>  
## 5 Wri      <tibble [3,627 × 26]>
```

What's going on here?

```
str(by_content$data)
```

```
## List of 5
## $ : tibble [3,627 × 26] (S3: tbl_df/tbl/data.frame)
##   ..$ ssid          : num [1:3627] 9466908 7683685 9025693 1009982
##   ..$ grade         : int [1:3627] 11 11 11 11 11 11 11 11 11 11 .
##   ..$ year          : int [1:3627] 11 11 11 11 11 11 11 11 11 11 .
##   ..$ testeventid   : num [1:3627] 148933 147875 143699 143962 150
##   ..$ asmtprmrydsblycd : num [1:3627] 0 10 40 82 10 80 50 10 50 82 ..
##   ..$ asmtscndrydsblycd : num [1:3627] 0 0 20 0 0 80 0 0 0 0 ...
##   ..$ Entry         : num [1:3627] 123 88 105 153 437 307 305 42 5
##   ..$ Theta         : num [1:3627] 1.27 1.55 3.28 4.48 2.67 ...
##   ..$ Status        : num [1:3627] 1 1 1 1 1 0 1 1 1 0 ...
##   ..$ Count         : num [1:3627] 36 36 36 36 36 36 36 36 36 36 .
##   ..$ RawScore      : num [1:3627] 23 25 33 35 31 36 34 18 3 36 ..
##   ..$ SE            : num [1:3627] 0.371 0.385 0.619 1.023 0.501 ..
##   ..$ Infit         : num [1:3627] 0.93 0.95 0.9 0.93 0.92 1 1.06
##   ..$ Infit_Z       : num [1:3627] -0.34 -0.37 -0.04 0.23 -0.18 0
##   ..$ Outfit        : num [1:3627] 0.82 0.81 1.63 0.35 0.88 1 0.86
##   ..$ Outfit_Z      : num [1:3627] -0.62 -0.56 1.03 -0.16 -0.12 0
##   ..$ Displacement  : num [1:3627] 0.0018 0.0019 0.0022 0.0023 0.0
##   ..$ PointMeasureCorr : num [1:3627] 0.42 0.42 0.3 0.27 0.31 0 0.14
##   ..$ Weight        : num [1:3627] 1 1 1 1 1 1 1 1 1 1 ...
##   ..$ ObservMatch    : num [1:3627] 75 80.6 91.7 97.2 86.1 100 94.4
##   ..$ ExpectMatch    : num [1:3627] 68.3 72 91.7 97.2 86.1 100 94.4
##   ..$ PointMeasureExpected: num [1:3627] 0.35 0.33 0.2 0.12 0.25 0.17
```

Explore a bit

```
map_dbl(by_content$data, nrow)
```

```
## [1] 3627 3629 3627 1435 3627
```

```
map_dbl(by_content$data, ncol)
```

```
## [1] 26 26 26 26 26
```

```
map_dbl(by_content$data, ~mean(.x$Theta))
```

```
## [1] 1.28001056 -0.06683086 1.37068376 1.57850321 1.26090709
```

It's a data frame!

We can add these summaries if we want

```
by_content %>%  
  mutate(n = map_dbl(data, nrow))
```

```
## # A tibble: 5 × 3  
## # Groups:   content [5]  
##   content data                                n  
##   <chr>    <list>                                <dbl>  
## 1 ELA      <tibble [3,627 × 26]>    3627  
## 2 Math     <tibble [3,629 × 26]>    3629  
## 3 Rdg      <tibble [3,627 × 26]>    3627  
## 4 Science  <tibble [1,435 × 26]>    1435  
## 5 Wri      <tibble [3,627 × 26]>    3627
```

map_*

- Note on the previous example we used `map_dbl` and we got a vector in return.
- What would happen if we just used `map`?

```
by_content %>%  
  mutate(n = map(data, nrow))
```

```
## # A tibble: 5 × 3  
## # Groups:   content [5]  
##   content data                                n  
##   <chr>    <list>                                <list>  
## 1 ELA     <tibble [3,627 × 26]> <int [1]>  
## 2 Math    <tibble [3,629 × 26]> <int [1]>  
## 3 Rdg     <tibble [3,627 × 26]> <int [1]>  
## 4 Science <tibble [1,435 × 26]> <int [1]>  
## 5 Wri     <tibble [3,627 × 26]> <int [1]>
```

Let's fit a model!

```
by_content %>%  
  mutate(m1 = map(data, ~lm(Theta ~ primary, data = .x)))
```

```
## # A tibble: 5 × 3  
## # Groups:   content [5]  
##   content data                                m1  
##   <chr>    <list>                                <list>  
## 1 ELA     <tibble [3,627 × 26]> <lm>  
## 2 Math    <tibble [3,629 × 26]> <lm>  
## 3 Rdg     <tibble [3,627 × 26]> <lm>  
## 4 Science <tibble [1,435 × 26]> <lm>  
## 5 Wri     <tibble [3,627 × 26]> <lm>
```


Extract the coefficients

```
by_content %>%  
  mutate(  
    m1 = map(data, ~lm(Theta ~ primary, data = .x)),  
    coefs = map(m1, coef)  
  )
```

```
## # A tibble: 5 × 4  
## # Groups:   content [5]  
##   content data          m1      coefs  
##   <chr>   <list>         <list> <list>  
## 1 ELA    <tibble [3,627 × 26]> <lm>   <dbl [11]>  
## 2 Math   <tibble [3,629 × 26]> <lm>   <dbl [12]>  
## 3 Rdg    <tibble [3,627 × 26]> <lm>   <dbl [11]>  
## 4 Science <tibble [1,435 × 26]> <lm>   <dbl [12]>  
## 5 Wri    <tibble [3,627 × 26]> <lm>   <dbl [12]>
```

Challenge

- Continue with the above, but output a data frame with three columns: **content**, **intercept**, and **TBI** (which is code 74).
- In other words, output the mean score for students who were coded as not having a disability (code 0), along with students coded as having TBI.

04:00

```
by_content %>%
  mutate(
    m1 = map(data, ~lm(Theta ~ primary, data = .x)),
    coefs = map(m1, coef),
    no_disab = map_dbl(coefs, 1),
    tbi = no_disab + map_dbl(coefs, "primary74")
  ) %>%
  select(content, no_disab, tbi)
```

```
## # A tibble: 5 × 3
## # Groups:   content [5]
##   content    no_disab      tbi
##   <chr>      <dbl>    <dbl>
## 1 ELA        0.9322336  0.1674462
## 2 Math      -0.1587907  0.1910821
## 3 Rdg        1.363101   1.629048
## 4 Science    1.491319   2.790971
## 5 Wri        1.571441   1.167429
```

Note – I wouldn't have necessarily expected you to add `no_disab` to the TBI coefficient.

Compare models

- Back to our original task – fit all three models

You try first

1. `lm(Theta ~ primary)`
2. `lm(Theta ~ primary + secondary)`
3. `lm(Theta ~ primary + secondary +
primary:secondary)`

04:00

Model fits

```
mods <- by_content %>%
  mutate(
    m1 = map(data, ~lm(Theta ~ primary, data = .x)),
    m2 = map(data, ~lm(Theta ~ primary + secondary, data = .x)),
    m3 = map(data, ~lm(Theta ~ primary * secondary, data = .x))
  )
mods
```

```
## # A tibble: 5 × 5
## # Groups:   content [5]
##   content data                m1      m2      m3
##   <chr>   <list>              <list> <list> <list>
## 1 ELA    <tibble [3,627 × 26]> <lm>   <lm>   <lm>
## 2 Math   <tibble [3,629 × 26]> <lm>   <lm>   <lm>
## 3 Rdg    <tibble [3,627 × 26]> <lm>   <lm>   <lm>
## 4 Science <tibble [1,435 × 26]> <lm>   <lm>   <lm>
## 5 Wri    <tibble [3,627 × 26]> <lm>   <lm>   <lm>
```

Brief foray into parallel iterations

The `stats::anova` function can compare the fit of two models

Pop Quiz

How would we extract just ELA model 1 and 2?

```
mods$m1[[1]]
```

```
##  
## Call:  
## lm(formula = Theta ~ primary, data = data_1)  
##  
## Coefficients:  
## (Intercept)      primary10  
##      0.93223      0.38570
```

```
mods$m2[[1]]
```

```
##  
## Call:  
## lm(formula = Theta ~ primary + second, data = data_2)  
##  
## Coefficients:  
## (Intercept)      primary10      primary20      primary40      primary50      primary20  
##      -0.03168      1.00143      0.42185      0.42185      0.42185      0.42185
```

Which fits better?

```
compare <- anova(mods$m1[[1]], mods$m2[[1]])  
compare
```

```
## Analysis of Variance Table  
##  
## Model 1: Theta ~ primary  
## Model 2: Theta ~ primary + secondary  
##   Res.Df    RSS Df Sum of Sq      F    Pr(>F)  
## 1     3616 20905  
## 2     3605 20100 11     804.26 13.113 < 2.2e-16 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

map2

- Works the same as `map` but iterates over two vectors concurrently
- Let's compare model 1 and 2

```
mods %>%  
  mutate(comp12 = map2(m1, m2, anova))
```

```
## # A tibble: 5 × 6  
## # Groups:   content [5]  
##   content data          m1      m2      m3      comp12  
##   <chr>   <list>      <list> <list> <list> <list>  
## 1 ELA     <tibble [3,627 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]>  
## 2 Math    <tibble [3,629 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]>  
## 3 Rdg     <tibble [3,627 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]>  
## 4 Science <tibble [1,435 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]>  
## 5 Wri     <tibble [3,627 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]>
```

Perhaps not terrifically helpful

Back to our **anova** object

- Can we pull out useful things?

```
str(compare)
```

```
## Classes 'anova' and 'data.frame':    2 obs. of  6 variables:
## $ Res.Df    : num  3616 3605
## $ RSS       : num  20905 20100
## $ Df        : num  NA 11
## $ Sum of Sq: num  NA 804
## $ F         : num  NA 13.1
## $ Pr(>F)    : num  NA 7.66e-25
## - attr(*, "heading")= chr [1:2] "Analysis of Variance Table\n" "Model 1"
```

Try pulling out the *p* value

Extract *p* value

- *Note – I'd recommend looking at more than just a p-value, but I do think this is useful for a quick glance*

```
compare$`Pr(>F)`
```

```
## [1] NA 7.663566e-25
```

```
compare[["Pr(>F)"]]
```

```
## [1] NA 7.663566e-25
```

```
compare$`Pr(>F)`[2]
```

```
## [1] 7.663566e-25
```

```
compare[["Pr(>F)"]][2]
```

```
## [1] 7.663566e-25
```

All p-values

Note – this is probably the most compact syntax, but that doesn't mean it's the most clear

```
mods %>%  
  mutate(comp12 = map2(m1, m2, anova),  
         p12 = map_dbl(comp12, list("Pr(>F)", 2)))
```

```
## # A tibble: 5 × 7  
## # Groups:   content [5]  
##   content data          m1      m2      m3      comp12  
##   <chr>   <list>      <list> <list> <list> <list>  
## 1 ELA    <tibble [3,627 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]> 7.6  
## 2 Math   <tibble [3,629 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]> 1.7  
## 3 Rdg    <tibble [3,627 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]> 1.5  
## 4 Science <tibble [1,435 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]> 4.6  
## 5 Wri    <tibble [3,627 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]> 5.7
```

Slight alternative

- Write a function that pulls the p-value from model comparison objects

```
extract_p <- function(anova_ob) {  
  anova_ob[["Pr(>F)"]][2]  
}
```

- Loop this function through the anova objects

```

mods %>%
  mutate(comp12 = map2(m1, m2, anova),
         p12 = map_dbl(comp12, extract_p))

```

```

## # A tibble: 5 × 7
## # Groups:   content [5]
##   content data          m1      m2      m3      comp12
##   <chr>   <list>         <list> <list> <list> <list>
## 1 ELA     <tibble [3,627 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]> 7.6
## 2 Math    <tibble [3,629 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]> 1.7
## 3 Rdg     <tibble [3,627 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]> 1.5
## 4 Science <tibble [1,435 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]> 4.6
## 5 Wri     <tibble [3,627 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]> 5.7

```

Brief sidetrack

We can also create the function using `purrr::compose()`.

Example

Create a centering function (which subtracts the mean from each ob)

```
center <- compose(~.x - mean(.x, na.rm = TRUE))
```

Use `~` and `.x`, just like with the `map()` functions.

Test it out

```
library(palmerpenguins)
```

```
penguins$bill_length_mm %>%  
  head()
```

```
## [1] 39.1 39.5 40.3    NA 36.7 39.3
```

```
penguins$bill_length_mm %>%  
  center() %>%  
  head()
```

```
## [1] -4.82193 -4.42193 -3.62193      NA -7.22193 -4.62193
```

```
penguins$bill_length_mm %>%  
  center() %>%  
  mean(na.rm = TRUE) %>%  
  round()
```

```
## [1] 0
```

Compose a p-val extractor

```
p <- compose(~.x[["Pr(>F)"]][2])
```

Use this instead

```
mods %>%  
  mutate(  
    comp12 = map2(m1, m2, anova),  
    p12 = map_dbl(comp12, p)  
  )
```

```
## # A tibble: 5 × 7  
## # Groups:   content [5]  
##   content data          m1      m2      m3      comp12  
##   <chr>   <list>      <list> <list> <list> <list>  
## 1 ELA    <tibble [3,627 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]> 7.6  
## 2 Math   <tibble [3,629 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]> 1.7  
## 3 Rdg    <tibble [3,627 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]> 1.5  
## 4 Science <tibble [1,435 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]> 4.6  
## 5 Wri    <tibble [3,627 × 26]> <lm>   <lm>   <lm>   <anova [2 × 6]> 5.7
```


Functions

This was a quick intro – don't worry if it doesn't really make sense yet. We'll talk about them (a lot) more in the coming weeks.

An alternative

Conducting operations by row

Operations by row

The `dplyr::rowwise()` function fundamentally changes the way a `tibble()` behaves

```
df <- tibble(name = c("Me", "You"), x = 1:2, y = 3:4, z = 5:6)
```

```
df %>%  
  mutate(m = mean(c(x, y, z)))
```

```
## # A tibble: 2 × 5  
##   name      x      y      z      m  
##   <chr> <int> <int> <int> <dbl>  
## 1 Me         1         3         5     3.5  
## 2 You         2         4         6     3.5
```

```
df %>%  
  rowwise() %>%  
  mutate(m = mean(c(x, y, z)))
```

```
## # A tibble: 2 × 5  
## # Rowwise:  
##   name      x      y      z      m  
##   <chr> <int> <int> <int> <dbl>  
## 1 Me         1         3         5     3  
## 2 You         2         4         6     4
```

Add a group & summarize

```
df %>%  
  rowwise(name) %>%  
  summarize(m = mean(c(x, y, z)))
```

```
## # A tibble: 2 × 2  
## # Groups:   name [2]  
##   name      m  
##   <chr> <dbl>  
## 1 Me      3  
## 2 You      4
```

List columns

If you apply rowwise operation with a list column, you don't have to loop

```
df <- tibble(var = list(1, 2:3, 4:6))
```

```
df %>%  
  mutate(  
    lngth = map_int(var, length)  
  )
```

```
## # A tibble: 3 × 2  
##   var      lngth  
##   <list>   <int>  
## 1 <dbl [1]>     1  
## 2 <int [2]>     2  
## 3 <int [3]>     3
```

```
df %>%  
  rowwise() %>%  
  mutate(lngth = length(var))
```

```
## # A tibble: 3 × 2  
## # Rowwise:  
##   var      lngth  
##   <list>   <int>  
## 1 <dbl [1]>     1  
## 2 <int [2]>     2  
## 3 <int [3]>     3
```

Creating list columns

You can use the `dplyr::nest_by()` function to create a list column for each group, *and* convert it to a rowwise data frame.

```
d %>%  
  nest_by(content)
```

```
## # A tibble: 5 × 2  
## # Rowwise:  content  
##   content          data  
##   <chr>    <list<tibble[,26]>>  
## 1 ELA      [3,627 × 26]  
## 2 Math     [3,629 × 26]  
## 3 Rdg      [3,627 × 26]  
## 4 Science  [1,435 × 26]  
## 5 Wri      [3,627 × 26]
```

Challenge

Given what we just learned, can you fit a model of the form **Theta ~ primary** to each content area (i.e., *not* using **{purrr}**)?

Wrap it in **list()** (should suggest this in the error reporting if you don't)

```
d %>%  
  nest_by(content) %>%  
  mutate(m1 = list(lm(Theta ~ primary, data = data)))
```

```
## # A tibble: 5 × 3  
## # Rowwise:   content  
##   content      data m1  
##   <chr>    <list<tibble[,26]>> <list>  
## 1 ELA      [3,627 × 26] <lm>  
## 2 Math     [3,629 × 26] <lm>  
## 3 Rdg      [3,627 × 26] <lm>  
## 4 Science  [1,435 × 26] <lm>  
## 5 Wri      [3,627 × 26] <lm>
```

02:00

Challenge 2

Can you extend it further and extract the coefficients with **coef**? What about creating a new column that has the intercept values?

```
d %>%  
  nest_by(content) %>%  
  mutate(m1 = list(lm(Theta ~ primary, data = data)),  
         coefs = list(coef(m1)))
```

```
## # A tibble: 5 × 4  
## # Rowwise:  content  
##   content                data m1      coefs  
##   <chr>    <list<tibble[,26]>> <list> <list>  
## 1 ELA      [3,627 × 26] <lm>    <dbl [11]>  
## 2 Math     [3,629 × 26] <lm>    <dbl [12]>  
## 3 Rdg      [3,627 × 26] <lm>    <dbl [11]>  
## 4 Science [1,435 × 26] <lm>    <dbl [12]>  
## 5 Wri      [3,627 × 26] <lm>    <dbl [12]>
```

02:00

Return atomic vectors

```
d %>%  
  nest_by(content) %>%  
  mutate(m1 = list(lm(Theta ~ primary, data = data)),  
         intercept = coef(m1)[1])
```

```
## # A tibble: 5 × 4  
## # Rowwise:   content  
##   content          data m1      intercept  
##   <chr>    <list<tibble[,26]>> <list>      <dbl>  
## 1 ELA      [3,627 × 26] <lm>      0.9322336  
## 2 Math     [3,629 × 26] <lm>     -0.1587907  
## 3 Rdg      [3,627 × 26] <lm>      1.363101  
## 4 Science  [1,435 × 26] <lm>      1.491319  
## 5 Wri      [3,627 × 26] <lm>      1.571441
```

Fit all models

The below gets us the same results we got before

```
mods2 <- d %>%
  nest_by(content) %>%
  mutate(
    m1 = list(lm(Theta ~ primary, data = data)),
    m2 = list(lm(Theta ~ primary + secondary, data = data)),
    m3 = list(lm(Theta ~ primary * secondary, data = data))
  )
mods2
```

```
## # A tibble: 5 × 5
## # Rowwise:  content
##   content      data m1      m2      m3
##   <chr>    <list<tibble[,26]>> <list> <list> <list>
## 1 ELA      [3,627 × 26] <lm>   <lm>   <lm>
## 2 Math     [3,629 × 26] <lm>   <lm>   <lm>
## 3 Rdg      [3,627 × 26] <lm>   <lm>   <lm>
## 4 Science  [1,435 × 26] <lm>   <lm>   <lm>
## 5 Wri      [3,627 × 26] <lm>   <lm>   <lm>
```

Look at all R^2

It's a normal data frame!

```
mods %>%
  pivot_longer(
    m1:m3,
    names_to = "model",
    values_to = "output"
  )
```

```
## # A tibble: 15 × 4
## # Groups:   content [5]
##   content data                model output
##   <chr>   <list>              <chr> <list>
## 1 ELA     <tibble [3,627 × 26]> m1     <lm>
## 2 ELA     <tibble [3,627 × 26]> m2     <lm>
## 3 ELA     <tibble [3,627 × 26]> m3     <lm>
## 4 Math    <tibble [3,629 × 26]> m1     <lm>
## 5 Math    <tibble [3,629 × 26]> m2     <lm>
## 6 Math    <tibble [3,629 × 26]> m3     <lm>
## 7 Rdg     <tibble [3,627 × 26]> m1     <lm>
## 8 Rdg     <tibble [3,627 × 26]> m2     <lm>
## 9 Rdg     <tibble [3,627 × 26]> m3     <lm>
## 10 Science <tibble [1,435 × 26]> m1     <lm>
## 11 Science <tibble [1,435 × 26]> m2     <lm>
```

Extract all R^2

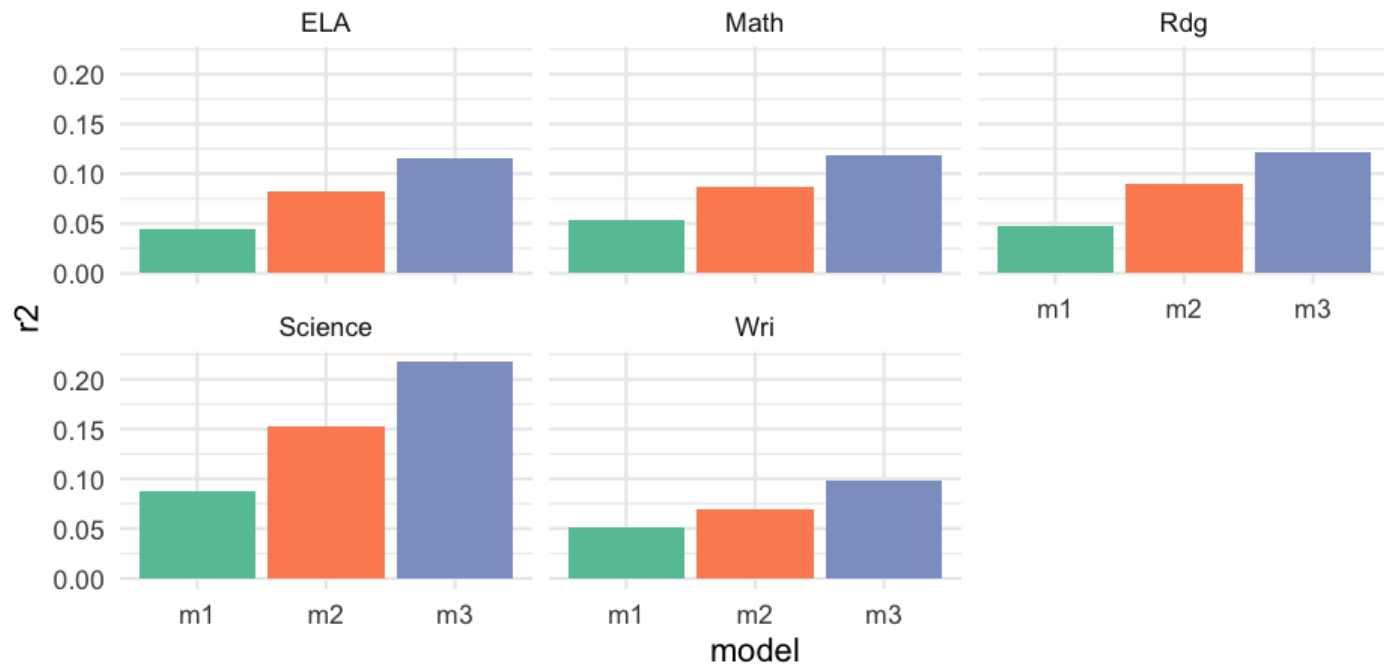
Note – might want to write a function here again

```
r2 <- mods %>%
  pivot_longer(
    m1:m3,
    names_to = "model",
    values_to = "output"
  ) %>%
  mutate(r2 = map_dbl(output, ~summary(.x)$r.squared))
r2
```

```
## # A tibble: 15 × 5
## # Groups:   content [5]
##   content data                model output      r2
##   <chr>   <list>              <chr> <list>    <dbl>
## 1 ELA     <tibble [3,627 × 26]> m1     <lm>    0.04517421
## 2 ELA     <tibble [3,627 × 26]> m2     <lm>    0.08190917
## 3 ELA     <tibble [3,627 × 26]> m3     <lm>    0.1161187
## 4 Math    <tibble [3,629 × 26]> m1     <lm>    0.05326550
## 5 Math    <tibble [3,629 × 26]> m2     <lm>    0.08675264
## 6 Math    <tibble [3,629 × 26]> m3     <lm>    0.1185931
## 7 Rdg     <tibble [3,627 × 26]> m1     <lm>    0.04805713
## 8 Rdg     <tibble [3,627 × 26]> m2     <lm>    0.08926212
## 9 Rdg     <tibble [3,627 × 26]> m3     <lm>    0.1217497
```

Plot

```
ggplot(r2, aes(model, r2)) +  
  geom_col(aes(fill = model)) +  
  facet_wrap(~content) +  
  guides(fill = "none") +  
  scale_fill_brewer(palette = "Set2")
```



Unnesting

- Sometimes you just want to `unnest`
- Imagine we want to plot the coefficients by model... how?
- `broom::tidy()` => `tidyr::unnest()`

Tidy

```
mods %>%
  pivot_longer(
    m1:m3,
    names_to = "model",
    values_to = "output"
  ) %>%
  mutate(tidied = map(output, broom::tidy))
```

```
## # A tibble: 15 × 5
## # Groups:   content [5]
##   content data          model output tidied
##   <chr>   <list>         <chr> <list> <list>
## 1 ELA     <tibble [3,627 × 26]> m1     <lm>   <tibble [11 × 5]>
## 2 ELA     <tibble [3,627 × 26]> m2     <lm>   <tibble [22 × 5]>
## 3 ELA     <tibble [3,627 × 26]> m3     <lm>   <tibble [132 × 5]>
## 4 Math    <tibble [3,629 × 26]> m1     <lm>   <tibble [12 × 5]>
## 5 Math    <tibble [3,629 × 26]> m2     <lm>   <tibble [23 × 5]>
## 6 Math    <tibble [3,629 × 26]> m3     <lm>   <tibble [144 × 5]>
## 7 Rdg     <tibble [3,627 × 26]> m1     <lm>   <tibble [11 × 5]>
## 8 Rdg     <tibble [3,627 × 26]> m2     <lm>   <tibble [22 × 5]>
## 9 Rdg     <tibble [3,627 × 26]> m3     <lm>   <tibble [132 × 5]>
## 10 Science <tibble [1,435 × 26]> m1     <lm>   <tibble [12 × 5]>
## 11 Science <tibble [1,435 × 26]> m2     <lm>   <tibble [22 × 5]>
## 12 Science <tibble [1,435 × 26]> m3     <lm>   <tibble [132 × 5]>
## 13 Wri    <tibble [3,627 × 26]> m1     <lm>   <tibble [12 × 5]>
```

Equivalently

```
mods %>%
  pivot_longer(
    m1:m3,
    names_to = "model",
    values_to = "output"
  ) %>%
  rowwise() %>%
  mutate(tidied = list(broom::tidy(output)))
```

```
## # A tibble: 15 × 5
## # Rowwise:  content
##   content data          model output tidied
##   <chr>   <list>         <chr> <list> <list>
## 1 ELA     <tibble [3,627 × 26]> m1     <lm>   <tibble [11 × 5]>
## 2 ELA     <tibble [3,627 × 26]> m2     <lm>   <tibble [22 × 5]>
## 3 ELA     <tibble [3,627 × 26]> m3     <lm>   <tibble [132 × 5]>
## 4 Math    <tibble [3,629 × 26]> m1     <lm>   <tibble [12 × 5]>
## 5 Math    <tibble [3,629 × 26]> m2     <lm>   <tibble [23 × 5]>
## 6 Math    <tibble [3,629 × 26]> m3     <lm>   <tibble [144 × 5]>
## 7 Rdg     <tibble [3,627 × 26]> m1     <lm>   <tibble [11 × 5]>
## 8 Rdg     <tibble [3,627 × 26]> m2     <lm>   <tibble [22 × 5]>
## 9 Rdg     <tibble [3,627 × 26]> m3     <lm>   <tibble [132 × 5]>
## 10 Science <tibble [1,435 × 26]> m1     <lm>   <tibble [12 × 5]>
## 11 Science <tibble [1,435 × 26]> m2     <lm>   <tibble [22 × 5]>
## 12 Science <tibble [1,435 × 26]> m3     <lm>   <tibble [132 × 5]>
```


Select and unnest

```
tidied <- mods %>%  
  gather(model, output, m1:m3) %>%  
  mutate(tidied = map(output, broom::tidy)) %>%  
  select(content, model, tidied) %>%  
  unnest(tidied)  
tidied
```


```
## # A tibble: 841 × 7  
## # Groups:   content [5]  
##   content model term          estimate std.error statistic    p.val  
##   <chr>    <chr> <chr>          <dbl>      <dbl>      <dbl>      <dbl>  
## 1 ELA      m1      (Intercept)  0.9322336  0.2150561  4.334839  1.498396e-04  
## 2 ELA      m1      primary10    0.3856986  0.2242965  1.719593  8.559207e-01  
## 3 ELA      m1      primary20   -0.03167527 0.7266436 -0.04359120 9.652327e-01  
## 4 ELA      m1      primary40   -1.844343   0.5559031 -3.317741  9.164595e-01  
## 5 ELA      m1      primary50    1.173722   0.2890447  4.060694  4.996391e-01  
## 6 ELA      m1      primary60    0.7762539   0.3866313  2.007737  4.474555e-01  
## 7 ELA      m1      primary70   -1.830257   0.3086128 -5.930595  3.301860e-01  
## 8 ELA      m1      primary74   -0.7647874  0.5182670 -1.475663  1.401215e-01  
## 9 ELA      m1      primary80    0.4676481  0.2428640  1.925556  5.423822e-01  
## 10 ELA     m1      primary82    0.3382547  0.2267600  1.491686  1.358687e-01  
## # ... with 831 more rows
```

Plot

Lets look how the primary coefficients change

```
to_plot <- names(coef(mods$m1[[1]]))

tidied %>%
  filter(term %in% to_plot) %>%
  ggplot(aes(estimate, term, color = model)) +
  geom_point() +
  scale_color_brewer(palette = "Set2") +
  facet_wrap(~content)
```



Last bit

- We've kind of been running the wrong models this whole time
- We forgot about grade!
- No problem, just change the grouping factor

By grade

```
by_grade_content <- d %>%  
  group_by(content, grade) %>%  
  nest()  
by_grade_content
```

```
## # A tibble: 31 × 3  
## # Groups:   content, grade [31]  
##   grade content data  
##   <int> <chr>   <list>  
## 1     11 ELA     <tibble [453 × 25]>  
## 2     11 Math    <tibble [460 × 25]>  
## 3     11 Rdg     <tibble [453 × 25]>  
## 4     11 Science <tibble [438 × 25]>  
## 5     11 Wri     <tibble [453 × 25]>  
## 6       3 ELA     <tibble [540 × 25]>  
## 7       3 Math    <tibble [536 × 25]>  
## 8       3 Rdg     <tibble [540 × 25]>  
## 9       3 Wri     <tibble [540 × 25]>  
## 10      4 ELA     <tibble [585 × 25]>  
## # ... with 21 more rows
```

Fit models

```
mods_grade <- by_grade_content %>%
  mutate(
    m1 = map(data, ~lm(Theta ~ primary, data = .x)),
    m2 = map(data, ~lm(Theta ~ primary + secondary,
                      data = .x)),
    m3 = map(data, ~lm(Theta ~ primary * secondary,
                      data = .x))
  )
mods_grade
```

```
## # A tibble: 31 × 6
## # Groups:   content, grade [31]
##   grade content data          m1      m2      m3
##   <int> <chr>   <list>    <list> <list> <list>
## 1     11 ELA    <tibble [453 × 25]> <lm>   <lm>   <lm>
## 2     11 Math   <tibble [460 × 25]> <lm>   <lm>   <lm>
## 3     11 Rdg    <tibble [453 × 25]> <lm>   <lm>   <lm>
## 4     11 Science <tibble [438 × 25]> <lm>   <lm>   <lm>
## 5     11 Wri    <tibble [453 × 25]> <lm>   <lm>   <lm>
## 6      3 ELA    <tibble [540 × 25]> <lm>   <lm>   <lm>
## 7      3 Math   <tibble [536 × 25]> <lm>   <lm>   <lm>
## 8      3 Rdg    <tibble [540 × 25]> <lm>   <lm>   <lm>
## 9      3 Wri    <tibble [540 × 25]> <lm>   <lm>   <lm>
## 10     4 ELA    <tibble [585 × 25]> <lm>   <lm>   <lm>
```

Look at R^2

```
mods_grade %>%
  pivot_longer(
    m1:m3,
    names_to = "model",
    values_to = "output"
  ) %>%
  mutate(r2 = map_dbl(output, ~summary(.x)$r.squared))
```

```
## # A tibble: 93 × 6
## # Groups:   content, grade [31]
##   grade content data          model output          r2
##   <int> <chr>   <list>          <chr> <list>          <dbl>
## 1     11 ELA    <tibble [453 × 25]> m1    <lm>          0.03353818
## 2     11 ELA    <tibble [453 × 25]> m2    <lm>          0.1084394
## 3     11 ELA    <tibble [453 × 25]> m3    <lm>          0.1536891
## 4     11 Math   <tibble [460 × 25]> m1    <lm>          0.1886003
## 5     11 Math   <tibble [460 × 25]> m2    <lm>          0.3161226
## 6     11 Math   <tibble [460 × 25]> m3    <lm>          0.4046634
## 7     11 Rdg    <tibble [453 × 25]> m1    <lm>          0.02066316
## 8     11 Rdg    <tibble [453 × 25]> m2    <lm>          0.1820512
## 9     11 Rdg    <tibble [453 × 25]> m3    <lm>          0.2337721
## 10    11 Science <tibble [438 × 25]> m1    <lm>          0.1259080
## # ... with 83 more rows
```

Plot

```
mods_grade %>%
  pivot_longer(
    m1:m3,
    names_to = "model",
    values_to = "output"
  ) %>%
  mutate(r2 = map_dbl(output, ~summary(.x)$r.squared)) %>%
  ggplot(aes(model, r2)) +
  geom_col(aes(fill = model)) +
  facet_grid(grade ~ content) +
  guides(fill = "none") +
  scale_fill_brewer(palette = "Set2")
```


Summary

- List columns are really powerful and really flexible
- Also help you stay organized
- You can approach the problem either with `{purrr}` or `dplyr::rowwise()`.
 - **Important:** If you use `rowwise()`, remember to `ungroup()` when you want it to go back to being a normal data frame
 - I'm asking you to learn both – the row-wise approach might be a bit easier but is a little less general (only works with data frames)

In-class

Midterm

Next time: Parallel iterations